

Sure! Let's dive into the topic of **Supervised Learning** and the subtopic of **Classification**.

## ✓ Supervised Learning

Supervised learning is one of the most common types of machine learning, where the model is trained on a labeled dataset. This means the algorithm learns from a set of input-output pairs, where the **input** data is provided with the **correct** answer (label). The goal of supervised learning is to learn a mapping from inputs to outputs, so the model can make predictions or classify new, unseen data.

### Key Characteristics:

- **Labeled Data:** The training data consists of inputs and their corresponding correct outputs (labels).
- **Objective:** The model learns the relationship between the input and the output so that it can predict the output for new, unseen inputs.
- **Types of Problems:** Supervised learning is generally used for regression (predicting continuous values) and classification (predicting discrete labels or categories).

## Classification (Subtopic)

**Classification** is a type of supervised learning where the output (or label) is a **category** or a **class**. Instead of predicting continuous values (like in regression), classification aims to assign each input to one of the predefined classes.

### Key Concepts in Classification:

1. **Classes:** The predefined categories or labels that the algorithm is trying to predict. For example, in an email spam filter, the classes could be "spam" and "not spam."
2. **Features:** The input data used to make predictions. For example, for email spam classification, features might include the email's subject, sender, and content.
3. **Model:** The algorithm or method used to classify the data. It learns the relationship between features (inputs) and the class labels (outputs).

## Classification Process:

- **Training Phase:** The model is provided with a labeled dataset, where the features of each example are paired with the correct class label. The model tries to learn the underlying patterns in the data.
- **Prediction Phase:** Once trained, the model can classify new, unseen data by predicting the class label for each new input based on the patterns it has learned.

## Common Classification Algorithms:

- **Logistic Regression:** Despite its name, it's used for binary classification (e.g., spam vs. not spam). It outputs probabilities that are then mapped to class labels.
- **Decision Trees:** A model that splits data into different branches based on feature values, resulting in leaf nodes representing class labels.
- **K-Nearest Neighbors (KNN):** Classifies data based on the majority class among the nearest neighbors of a data point.
- **Support Vector Machines (SVM):** Finds a hyperplane that best separates the data into different classes.
- **Naive Bayes:** A probabilistic model based on Bayes' theorem, often used for text classification tasks.
- **Random Forest:** An ensemble of decision trees that helps reduce overfitting and improve prediction accuracy.

**Example of a Classification Problem:** Let's say you're tasked with predicting whether a patient has a disease based on medical data (e.g., age, blood pressure, and cholesterol levels). The output (or class) could be "**disease**" or "**no disease**". The model would be trained on historical medical data with the correct labels, and it would learn to classify new patients as either having the disease or not.

## Evaluation Metrics for Classification

To assess the performance of classification models, we use several evaluation metrics:

1. **Accuracy:** The percentage of correct predictions  $\$ \text{ (total correct predictions / total predictions)} \$$ .

2. **Precision:** The proportion of positive predictions that are actually correct \$  $(\text{true positives} / (\text{true positives} + \text{false positives}))$ .
3. **Recall (Sensitivity):** The proportion of actual positive cases correctly identified \$  $(\text{true positives} / (\text{true positives} + \text{false negatives}))$ .
4. **F1-Score:** The harmonic mean of precision and recall, giving a balanced measure \$  $(2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall}))$ .
5. **Confusion Matrix:** A table showing the true positives, true negatives, false positives, and false negatives.

## Example Scenario:

Imagine you're using a classification model to detect whether a user's online behavior is legitimate or fraudulent. You have historical data with labeled examples of "fraudulent" and "legitimate" behaviors. The model would be trained on this labeled data and would then predict whether new behavior is "fraudulent" or "legitimate."

## Summary of Classification Process:

1. **Collect and preprocess data:** Gather labeled data (features and corresponding labels).
2. **Choose a classification algorithm:** Select an appropriate algorithm (e.g., decision trees, SVM, or logistic regression).
3. **Train the model:** Feed the data into the algorithm so it can learn the mapping between features and labels.
4. **Evaluate the model:** Test the model on new data and evaluate its performance using metrics like accuracy, precision, recall, and F1-score.
5. **Deploy the model:** Once satisfied with the model's performance, deploy it to classify new, unseen data.

## Real-World Applications of Classification:


- **Medical Diagnostics:** Classifying medical images or test results as healthy or diseased.
- **Email Filtering:** Classifying emails as spam or not spam.
- **Credit Scoring:** Classifying loan applicants as creditworthy or not.
- **Sentiment Analysis:** Classifying text as positive, negative, or neutral.

If you need further clarification or examples on specific points, feel free to ask!

**Types of Supervised Learning Problems:** 1. **Classification:** This involves predicting a categorical outcome. For instance, a classification algorithm could determine whether an email is spam or not based on its content. 2. **Regression:** This deals with predicting a continuous numerical value. For example, using regression algorithms to predict house prices based on features like size, location, and amenities.

**Classification in Machine Learning** Classification is the process of categorizing data into predefined classes or groups based on their features. In Supervised Learning, the model is trained on a labeled dataset and then used to predict the class or category of new, unseen data. For example, a classification model trained on images labeled as "dogs" or "cats" can predict the category of new images based on their features such as color or shape.

**Types of Classification:** 1. **Binary Classification:** The goal is to classify data into one of two categories. For example, predicting whether a person has a specific disease based on their health data. 2. **Multiclass Classification:** Here, the input is classified into one of several categories. For example, predicting the species of a flower based on its characteristics.



<https://www.researchgate.net/profile/Emerson-Nithiyaraj/publication/342987800/figure/fig1/AS:913942624870401@1594912310090/Binary-vs-Multiclass-classification.jpg>

### Other Classification Types:

- **Multi-Label Classification:** Each data point can belong to multiple categories. For example, an image could be classified as both "cat" and "dog" if it contains both animals.
- **Imbalanced Classification:** This type of classification deals with datasets where one class is heavily underrepresented. An example is predicting rare diseases, where the majority of the data belongs to common conditions, and the minority class represents the rare disease.

**Classification Algorithms** There are various types of classifiers algorithms. Some of them are :

**Linear Classifiers** Linear models create a linear decision boundary between classes. They are simple and computationally efficient. Some of the linear classification models are as follows: \* Logistic Regression \* Support Vector Machines having kernel = 'linear' \* Single-layer Perceptron \* Stochastic Gradient Descent (SGD) Classifier

### Non-linear Classifiers

Non-linear models create a non-linear decision boundary between classes. They can capture more complex relationships between the input features and the target variable.

## ✓ Evaluating a **classification model**

in machine learning is crucial to understanding its performance and ability to generalize to unseen data. Various metrics and techniques are used to assess models, depending on the nature of the problem. Here are key evaluation methods: 1.

**Classification Accuracy:** This is the proportion of correctly classified instances out of the total instances. It is easy to understand but can be misleading for imbalanced datasets where the majority class dominates the accuracy score. 2. **Confusion**

**Matrix:** This table displays true positives, true negatives, false positives, and false negatives, helping to calculate more detailed evaluation metrics. 3. **Precision and**

**Recall:** o Precision: The proportion of true positives over all predicted positives, indicating how many selected items were relevant. o Recall: The proportion of true positives over all actual positives, showing how well the model captures relevant items. These metrics are especially useful when dealing with imbalanced data or when false positives and false negatives have different costs. 4. **F1-Score:** The harmonic mean of precision and recall, balancing the two metrics, particularly useful for imbalanced datasets where both false positives and false negatives matter. 5.

**ROC Curve and AUC:** \* ROC Curve: A plot of the true positive rate (recall) against the false positive rate, showing performance at different thresholds. \* AUC: The Area Under the Curve measures overall performance, with values from 0.5 (random guessing) to 1 (perfect classification). 6. **Cross-Validation:** This technique divides the dataset into multiple subsets, training the model on some while testing on others. It provides a more reliable estimate of model performance and helps avoid overfitting.

Choosing the right metric(s) depends on the problem's specific requirements, and models should always be evaluated on independent test data to ensure robust performance.

```
# Importing the required libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn import datasets
from sklearn import svm
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.naive_bayes import GaussianNB

# import the iris dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

# splitting X and y into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=1)

# GAUSSIAN NAIVE BAYES
gnb = GaussianNB()
# train the model
gnb.fit(X_train, y_train)
# make predictions
gnb_pred = gnb.predict(X_test)
# print the accuracy
print("Accuracy of Gaussian Naive Bayes: ",
      accuracy_score(y_test, gnb_pred))
# print other performance metrics
print("Precision of Gaussian Naive Bayes: ",
      precision_score(y_test, gnb_pred, average='weighted'))
print("Recall of Gaussian Naive Bayes: ",
      recall_score(y_test, gnb_pred, average='weighted'))
print("F1-Score of Gaussian Naive Bayes: ",
      f1_score(y_test, gnb_pred, average='weighted'))

# DECISION TREE CLASSIFIER
dt = DecisionTreeClassifier(random_state=0)
# train the model
dt.fit(X_train, y_train)
# make predictions
dt_pred = dt.predict(X_test)
# print the accuracy
print("Accuracy of Decision Tree Classifier: ",
      accuracy_score(y_test, dt_pred))
# print other performance metrics
print("Precision of Decision Tree Classifier: ",
      precision_score(y_test, dt_pred, average='weighted'))
print("Recall of Decision Tree Classifier: ",
      recall_score(y_test, dt_pred, average='weighted'))
print("F1-Score of Decision Tree Classifier: ",
      f1_score(y_test, dt_pred, average='weighted'))
```

```
# SUPPORT VECTOR MACHINE
svm_clf = svm.SVC(kernel='linear') # Linear Kernel
# train the model
svm_clf.fit(X_train, y_train)
# make predictions
svm_clf_pred = svm_clf.predict(X_test)
# print the accuracy
print("Accuracy of Support Vector Machine: ",
      accuracy_score(y_test, svm_clf_pred))
# print other performance metrics
print("Precision of Support Vector Machine: ",
      precision_score(y_test, svm_clf_pred, average='weighted'))
print("Recall of Support Vector Machine: ",
      recall_score(y_test, svm_clf_pred, average='weighted'))
print("F1-Score of Support Vector Machine: ",
      f1_score(y_test, svm_clf_pred, average='weighted'))
```

```
Accuracy of Gaussian Naive Bayes: 0.9333333333333333
Precision of Gaussian Naive Bayes: 0.9352007469654529
Recall of Gaussian Naive Bayes: 0.9333333333333333
F1-Score of Gaussian Naive Bayes: 0.933615520282187
Accuracy of Decision Tree Classifier: 0.9555555555555556
Precision of Decision Tree Classifier: 0.9555555555555556
Recall of Decision Tree Classifier: 0.9555555555555556
F1-Score of Decision Tree Classifier: 0.9555555555555556
Accuracy of Support Vector Machine: 1.0
Precision of Support Vector Machine: 1.0
Recall of Support Vector Machine: 1.0
F1-Score of Support Vector Machine: 1.0
```

## ✓ Supervised Learning

Supervised learning for classification involves building models to predict discrete labels or classes. Here's an in-depth guide on theoretical and practical aspects for learning about logistic regression, decision trees, evaluation metrics, and building classification models using Scikit-Learn.

### 1. Theoretical

#### 1.1. Logistic Regression:



**! Logistic regression is a classification algorithm !** used for binary or multi-class classification. Despite its name, it's a linear model for classification, not regression.

- **Core Concepts:**

- **Logistic Function (Sigmoid):** Used to map predicted values to probabilities. The sigmoid function outputs values between 0 and 1, allowing for classification. 
$$P(y=1|x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n)}}$$
 
$$P(y=1|x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n)}}$$
- **Binary Classification:** Logistic regression is used when the target variable has two classes (e.g., spam/not spam).
- **Multi-class Classification (One-vs-Rest):** Logistic regression can be extended to classify more than two classes using a strategy like "One-vs-Rest" (OvR).
- **Decision Boundary:** Logistic regression separates classes using a linear decision boundary, making it a linear classifier.

- **Mathematical Intuition:**

- **Cost Function:** Uses **log-loss** or binary cross-entropy to measure the error in predictions.
- **Optimization:** Typically optimized using gradient descent or advanced solvers like **L-BFGS**.

- **Further Reading:**

- **Books:**
  - "An Introduction to Statistical Learning" (Chapter 4) – Gareth James et al.
  - "Elements of Statistical Learning" (Chapter 4) – Hastie et al.
- **Scikit-Learn Documentation:**
  - [Logistic Regression](#).

## 1.2. Decision Trees:

Decision trees are non-linear classifiers that recursively split the data into subsets

based on feature values.

- **Core Concepts:**

- **Tree Structure:** Composed of nodes (decision points) and branches (outcomes), starting from a root and splitting down to leaves (class labels).
- **Splitting Criteria:** Common splitting criteria include:
  - **Gini Index:** Measures impurity; a pure node has a Gini index of 0.
  - **Entropy/Information Gain:** Measures reduction in uncertainty.
  - **Chi-Square Test:** Used to evaluate the statistical significance of splits.
- **Pruning:** Reduces overfitting by trimming unnecessary branches.
- **Advantages:** Interpretable, non-linear, and works well with both categorical and numerical data.

- **Mathematical Intuition:**

- **Gini Impurity Formula:**  $Gini = 1 - \sum (p_i^2)$  where  $(p_i)$  is the probability of class  $(i)$  at a node.

- **Further Reading:**

- **Books:**
  - "Pattern Recognition and Machine Learning" by Christopher Bishop.
  - "Data Mining: Practical Machine Learning Tools" by Ian Witten et al.
- **Articles:**
  - [Decision Trees Explained](#) on Towards Data Science.
- **Scikit-Learn Documentation:**
  - [Decision Trees](#).

## 1.3. K-Nearest Neighbors (K-NN):

K-Nearest Neighbors (K-NN) is a non-parametric, instance-based learning algorithm used for both classification and regression. It predicts outcomes by identifying the  $(k)$ -nearest data points in the feature space based on a similarity measure.

# Core Concepts:

## 1. Instance-Based Learning:

- Unlike parametric models, K-NN does not assume an underlying functional form for the data. Instead, it memorizes the training data and makes predictions based on the nearest neighbors.

## 2. Distance Metrics:

- Used to measure the similarity between points. Common metrics include:

- **Euclidean Distance:** For continuous variables.

$$d(p, q) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

$$d(p, q) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$

- **Manhattan Distance:** For absolute differences.

$$d(p, q) = \sum_{i=1}^n |p_i - q_i|$$

$$d(p, q) = \sum_{i=1}^n |p_i - q_i|$$

- **Hamming Distance:** For categorical data, measures differing features.

## 3. Classification:

- The algorithm assigns a class label based on the majority class among the ( k )-nearest neighbors.

## 4. Regression:

- Predicts a continuous value by averaging the target values of the ( k )-nearest neighbors.

## 5. Choosing ( k ):

- A small ( k ): More sensitive to noise and leads to overfitting.
- A large ( k ): Produces smoother predictions but may underfit the data.

## 6. Decision Boundary:

- K-NN creates non-linear decision boundaries influenced by the distribution of training data.

---

# Mathematical Intuition:

### 1. Distance Calculation:

- The similarity between data points is measured using a distance metric (e.g., Euclidean, Manhattan).

### 2. Majority Voting (Classification):

- For a given data point (  $x$  ), the algorithm looks at the (  $k$  )-closest neighbors and assigns the class label with the highest frequency.

### 3. Averaging (Regression):

- For regression tasks, the algorithm computes the mean of the target values of the (  $k$  )-nearest neighbors:

$$\hat{y} = \frac{1}{k} \sum_{i=1}^k y_i$$

### 4. Effect of Feature Scaling:

- Distance calculations are sensitive to feature magnitudes, so normalizing or standardizing features is crucial.

### 5. Complexity:

- **Training Time:** Minimal since no model is built.
- **Prediction Time:** (  $O(n)$  ), where (  $n$  ) is the size of the training set.

## Further Reading:

### Books:

1. "Introduction to Statistical Learning" – Gareth James et al. (Chapter 2).
2. "Pattern Recognition and Machine Learning" – Christopher Bishop (Chapter 8).

### Articles:

1. "Understanding K-Nearest Neighbors" – Towards Data Science.
2. "K-Nearest Neighbors in Machine Learning" – Analytics Vidhya.

### Scikit-Learn Documentation:

1. [K-Neighbors Classifier](#).

## 2. [K-Neighbors Regressor](#).

K-NN is an essential algorithm for beginners and provides a strong foundation for understanding instance-based learning and distance metrics in machine learning.

## 1.4. Supervised Learning: Naive Bayes:

Naive Bayes is a family of simple but effective probabilistic algorithms based on Bayes' Theorem, used primarily for classification tasks. It is particularly popular for text classification and problems with high-dimensional input data.

### Core Concepts

#### 1. Bayes' Theorem:

- The algorithm uses Bayes' Theorem to calculate the probability of a class given the features:  

$$P(C|X) = \frac{P(X|C) \cdot P(C)}{P(X)}$$

$$P(C|X) = \frac{P(X|C) \cdot P(C)}{P(X)}$$
 Where:
  - ( $P(C|X)$ ): Posterior probability of class ( $C$ ) given features ( $X$ ).
  - ( $P(X|C)$ ): Likelihood of features ( $X$ ) given class ( $C$ ).
  - ( $P(C)$ ): Prior probability of class ( $C$ ).
  - ( $P(X)$ ): Marginal probability of features ( $X$ ).

#### 2. Naive Assumption:

- The algorithm assumes that all features are **conditionally independent** given the class label, which simplifies the computation of ( $P(X|C)$ ) as:
 
$$P(X|C) = P(x_1|C) \cdot P(x_2|C) \cdot \dots \cdot P(x_n|C)$$

$$P(X|C) = P(x_1|C) \cdot P(x_2|C) \cdot \dots \cdot P(x_n|C)$$
 This "naive" independence assumption is rarely true in practice but often works well.

#### 3. Types of Naive Bayes Classifiers:

- **Gaussian Naive Bayes:** Used for continuous data assuming a Gaussian

(normal) distribution.

- **Multinomial Naive Bayes:** Suitable for discrete data, often used in text classification tasks.
- **Bernoulli Naive Bayes:** Designed for binary/boolean data, useful for presence/absence features.

#### 4. Applications:

- Spam detection.
- Sentiment analysis.
- Document categorization.
- Medical diagnosis.

#### 5. Advantages:

- Simple and fast to train.
- Performs well with small datasets.
- Effective for high-dimensional input data.

#### 6. Limitations:

- Strong independence assumption may lead to inaccuracies if features are correlated.
- May struggle with datasets containing rare feature values or small training data.

## Mathematical Intuition

### 1. Posterior Probability:

- Naive Bayes predicts the class ( $\hat{C}$ ) that maximizes the posterior probability:  $\hat{C} = \underset{C}{\text{argmax}} \, P(C|X)$   
 $\hat{C} = \underset{C}{\text{argmax}} \, P(C|X)$  Using Bayes' Theorem:  $\hat{C} = \underset{C}{\text{argmax}} \, \frac{P(X|C) \cdot P(C)}{P(X)}$   
 $\hat{C} = \underset{C}{\text{argmax}} \, \frac{P(X|C) \cdot P(C)}{P(X)}$  Since ( $P(X)$ ) is constant for all classes, it simplifies to:  
 $\hat{C} = \underset{C}{\text{argmax}} \, P(X|C) \cdot P(C)$   
 $\hat{C} = \underset{C}{\text{argmax}} \, P(X|C) \cdot P(C)$

### 2. Class Likelihood:

- For continuous features, Gaussian Naive Bayes assumes a normal distribution: 
$$P(x_i|C) = \frac{1}{\sqrt{2\pi}\sigma_C} \exp\left(-\frac{(x_i - \mu_C)^2}{2\sigma_C^2}\right)$$
 
$$P(x_i|C) = \frac{1}{\sqrt{2\pi}\sigma_C} \exp\left(-\frac{(x_i - \mu_C)^2}{2\sigma_C^2}\right)$$
 Where  $(\mu_C)$  and  $(\sigma_C^2)$  are the mean and variance of the feature  $(x_i)$  for class  $(C)$ .

### 3. Logarithmic Simplification:

- To avoid underflow and simplify computation, Naive Bayes often uses the logarithm of probabilities: 
$$\log(P(C|X)) = \log(P(C)) + \sum_{i=1}^n \log(P(x_i|C))$$
 
$$\log(P(C|X)) = \log(P(C)) + \sum_{i=1}^n \log(P(x_i|C))$$

### 4. Feature Scaling:

- Naive Bayes is not sensitive to feature scaling, as it relies on probabilities rather than distances.

---

## Further Reading

### Books:

- "Introduction to Statistical Learning" – Gareth James et al.
  - Provides an excellent introduction to Naive Bayes with practical examples.
- "Pattern Recognition and Machine Learning" – Christopher Bishop.
  - Offers a deeper dive into Bayesian methods.

### Articles:

- "Understanding Naive Bayes" – Towards Data Science.
- "Text Classification Using Naive Bayes" – Analytics Vidhya.

### Scikit-Learn Documentation:

- [GaussianNB Classifier.](#)
- [MultinomialNB Classifier.](#)
- [BernoulliNB Classifier.](#)

Naive Bayes is an essential algorithm in the supervised learning toolkit, particularly

when working with probabilistic models and high-dimensional data. Its simplicity, efficiency, and ability to handle real-world problems make it a popular choice for beginners and experts alike.<sup>1</sup>

## 1.5. Supervised Learning: Support Vector Machines (SVM):

Support Vector Machines (SVM) are a class of supervised learning algorithms used for both classification and regression tasks. They are particularly effective in high-dimensional spaces and are known for their ability to find a robust decision boundary between classes.

---

### Core Concepts

#### 1. Key Idea:

- SVM aims to find the **optimal hyperplane** that separates classes with the **maximum margin** in a feature space.
- For non-linearly separable data, it uses **kernels** to project data into a higher-dimensional space where it becomes linearly separable.

#### 2. Support Vectors:

- The data points closest to the decision boundary (margin) are called **support vectors**. They are critical to defining the position of the hyperplane.

#### 3. Decision Boundary:

- SVM defines the decision boundary as a hyperplane. For two-dimensional data, this is a line.
- The equation for the hyperplane is: 
$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

Where  $(\mathbf{w})$  is the weight vector,  $(\mathbf{x})$  is the input feature vector, and  $(b)$  is the bias term.

#### 4. Hard Margin vs. Soft Margin:

- **Hard Margin:** Assumes perfect linear separability of classes.



- **Soft Margin:** Allows some misclassification to balance margin size and classification accuracy.

## 5. Kernels:

- Kernels allow SVM to operate in high-dimensional spaces by mapping data into a different feature space. Common kernels include:
  - **Linear Kernel:**  $K(\mathbf{x}, \mathbf{y}) = \mathbf{x} \cdot \mathbf{y}$ .
  - **Polynomial Kernel:**  $K(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y} + c)^d$ .
  - **Radial Basis Function (RBF) Kernel:**  $K(\mathbf{x}, \mathbf{y}) = \exp(-\gamma \|\mathbf{x} - \mathbf{y}\|^2)$ .
  - **Sigmoid Kernel:** Mimics a neural network activation function.

## 6. Applications:

- Text classification (e.g., spam detection).
- Image recognition.
- Bioinformatics (e.g., protein classification).

## 7. Advantages:

- Effective in high-dimensional spaces.
- Robust against overfitting when the number of features exceeds the number of samples.
- Can handle non-linear decision boundaries using kernels.

## 8. Limitations:

- Computationally expensive for large datasets.
- Requires careful parameter tuning (e.g., kernel choice,  $C$ ,  $\gamma$ ).
- Not suitable for datasets with significant overlap between classes.

---

# Mathematical Intuition

## 1. Optimization Problem:

- SVM maximizes the margin  $\frac{1}{\|\mathbf{w}\|}$  subject to constraints for all data points  $(i)$ :

$$y_i (\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$$

$$\mathbf{w} \cdot \mathbf{x}_i + b \geq 1$$

Where ( $y_i \in \{-1, 1\}$ ) represents the class labels.

## 2. Soft Margin Optimization:

- Introduces a slack variable ( $\xi_i$ ) to allow misclassification:
 
$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$
 Where ( $C$ ) is a regularization parameter that controls the trade-off between margin size and classification error.

## 3. Kernel Trick:

- Instead of computing the dot product in a high-dimensional space, kernels allow calculations in the original feature space:
 
$$K(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x}) \cdot \phi(\mathbf{y})$$

$$K(\mathbf{x}, \mathbf{y}) = \phi(\mathbf{x}) \cdot \phi(\mathbf{y})$$

Where  $\phi$  maps input data into a higher-dimensional feature space.

## 4. Dual Formulation:

- SVM often solves the **dual optimization problem** to efficiently handle kernels:
 
$$\max_{\alpha} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j)$$

$$\max_{\alpha} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j)$$
 Subject to  $(\sum_{i=1}^n \alpha_i y_i = 0)$  and  $(0 \leq \alpha_i \leq C)$ .

## 5. Prediction:

- The predicted class for a point ( $\mathbf{x}$ ) is:
 
$$f(\mathbf{x}) = \text{sign} \left( \sum_{i=1}^n \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b \right)$$

$$f(\mathbf{x}) = \text{sign} \left( \sum_{i=1}^n \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b \right)$$

## Further Reading

### Books:

1. **"An Introduction to Statistical Learning"** – Gareth James et al. (Chapter 9).
  - Provides an intuitive explanation with examples.
2. **"The Elements of Statistical Learning"** – Hastie, Tibshirani, Friedman (Chapter 12).
  - A more technical dive into SVM theory and extensions.
3. **"Pattern Recognition and Machine Learning"** – Christopher Bishop (Chapter 7).
  - Excellent for mathematical details.

## Articles:

1. "Understanding Support Vector Machines" – Towards Data Science.
2. "A Beginner's Guide to SVMs" – Analytics Vidhya.

## Scikit-Learn Documentation:

1. [SVC Classifier](#).
2. [SVR Regressor](#).

SVM is a powerful algorithm for both classification and regression, providing robust performance in many practical scenarios. Its theoretical foundation and versatility make it a staple in the supervised learning domain.

## 2. Theoretical: Evaluation Metrics for Classification Models

Once the model is trained, it's critical to evaluate its performance using relevant metrics. These metrics help to understand how well the classifier generalizes to unseen data.

### 2.1. Common Evaluation Metrics:

- **Accuracy:** The proportion of correct predictions.  $\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$
- **Precision:** The proportion of true positives among predicted positives.  $\text{Precision} = \frac{TP}{TP + FP}$
- **Recall (Sensitivity):** The proportion of actual positives that were correctly classified.  $\text{Recall} = \frac{TP}{TP + FN}$
- **F1-Score:** The harmonic mean of precision and recall, useful when dealing with imbalanced classes.  $F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$
- **Confusion Matrix:** A matrix used to evaluate the performance of a classification model. It shows the counts of true positives, true negatives, false positives, and false negatives.

### 2.2. ROC and AUC:

- **ROC Curve:** Plots the true positive rate (TPR) against the false positive rate (FPR) at various thresholds.
- **AUC (Area Under the Curve):** Measures the overall performance of the classifier. AUC = 1 indicates a perfect classifier, while AUC = 0.5 indicates random guessing.

## ✓ 3. Practical: Building Classification Models

### 3.1. Logistic Regression: Practical Example

This example demonstrates how to build a logistic regression model in Python, visualize the classification, and evaluate its performance using metrics.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, (

# Generate synthetic dataset
X, y = make_classification(n_samples=100, n_features=2, n_informative=

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.

# Logistic Regression Model
model = LogisticRegression()
model.fit(X_train, y_train)


# Predictions
y_pred = model.predict(X_test)

# Visualize decision boundary
def plot_decision_boundary(X, y, model):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                          np.arange(y_min, y_max, 0.01))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Pa
    plt.title("Logistic Regression Decision Boundary")
    plt.xlabel("Feature 1")
```


```
plt.ylabel("Feature 2")
plt.show()

plot_decision_boundary(X_test, y_test, model)

# Evaluation Metrics
print("Classification Report:\n", classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()
```



|                        |           |        |          |         |
|------------------------|-----------|--------|----------|---------|
| Classification Report: |           |        |          |         |
|                        | precision | recall | f1-score | support |
| 0                      | 1.00      | 0.94   | 0.97     | 16      |
| 1                      | 0.93      | 1.00   | 0.97     | 14      |
| accuracy               |           |        | 0.97     | 30      |
| macro avg              | 0.97      | 0.97   | 0.97     | 30      |
| weighted avg           | 0.97      | 0.97   | 0.97     | 30      |



# Explanation

## 1. Dataset Creation:

- A synthetic dataset with two features and two classes is generated using `make_classification`.

## 2. Model Training:

- A `LogisticRegression` model is trained using the training data.

## 3. Decision Boundary Visualization:

- The decision boundary of the model is plotted using `contourf`, showing how the model separates the classes.

## 4. Evaluation Metrics:

- `classification_report` provides precision, recall, F1-score, and accuracy for both classes.
- `confusion_matrix` and its visualization display the number of correct and incorrect predictions for each class.

---

## Key Metrics in Classification

- **Precision:**  $\frac{\text{TP}}{\text{TP} + \text{FP}}$
- **Recall:**  $\frac{\text{TP}}{\text{TP} + \text{FN}}$
- **F1-Score:** Harmonic mean of precision and recall.
- **Accuracy:**  $\frac{(\text{TP} + \text{TN})}{\text{Total}}$

This example gives a concise yet comprehensive overview of logistic regression in practice, with visual and numerical insights.

## 3.2. Decision Trees: Practical Example

This example demonstrates building a decision tree classifier in Python, visualizing the tree and classification boundaries, and evaluating its performance using metrics.



```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, (

# Generate synthetic dataset
X, y = make_classification(n_samples=100, n_features=2, n_informative=

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.

# Decision Tree Model
model = DecisionTreeClassifier(max_depth=3, random_state=42)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Visualize Decision Boundary
def plot_decision_boundary(X, y, model):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                          np.arange(y_min, y_max, 0.01))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Pa
    plt.title("Decision Tree Decision Boundary")
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 2")
    plt.show()

plot_decision_boundary(X_test, y_test, model)

# Visualize the Decision Tree
plt.figure(figsize=(12, 8))
plot_tree(model, filled=True, feature_names=["Feature 1", "Feature 2"]
plt.title("Decision Tree Visualization")
plt.show()


# Evaluation Metrics
print("Classification Report:\n", classification_report(y_test, y_pred

```


```

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()

```



Decision Tree Visualization



value = [31, 0]  
class = Class 0

value = [3, 3]  
class = Class 0


gini = 0.0  
samples = 3  
value = [0, 3]  
class = Class 1

gini = 0.0  
samples = 3  
value = [3, 0]  
class = Class 0

#### Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.94   | 0.97     | 16      |
| 1            | 0.93      | 1.00   | 0.97     | 14      |
| accuracy     |           |        | 0.97     | 30      |
| macro avg    | 0.97      | 0.97   | 0.97     | 30      |
| weighted avg | 0.97      | 0.97   | 0.97     | 30      |

#### Confusion Matrix



# Explanation

## 1. Dataset Creation:

- A synthetic dataset with two features and two classes is generated using `make_classification`.

## 2. Model Training:

- A `DecisionTreeClassifier` with a maximum depth of 3 is trained using the training data.

## 3. Decision Boundary Visualization:

- The decision boundary of the model is plotted using `contourf`, illustrating how the decision tree splits the feature space.

## 4. Tree Visualization:

- `plot_tree` is used to display the structure of the decision tree, showing splits, feature importance, and class predictions at each node.

## 5. Evaluation Metrics:

- `classification_report` provides precision, recall, F1-score, and accuracy for both classes.
- `confusion_matrix` and its visualization display the number of correct and incorrect predictions for each class.

---

## Key Metrics in Classification

- **Precision:**  $\frac{\text{TP}}{\text{TP} + \text{FP}}$
- **Recall:**  $\frac{\text{TP}}{\text{TP} + \text{FN}}$
- **F1-Score:** Harmonic mean of precision and recall.
- **Accuracy:**  $\frac{\text{TP} + \text{TN}}{\text{Total}}$

This concise example covers practical implementation, visualization, and evaluation of decision trees for classification.

### ✓ 3.3. K-Nearest Neighbors (K-NN): Practical Example

This example demonstrates using K-NN for classification in Python, visualizing the classification boundaries, and evaluating its performance.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, (

# Generate synthetic dataset
X, y = make_classification(n_samples=100, n_features=2, n_informative=

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.


# K-NN Model
model = KNeighborsClassifier(n_neighbors=5)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Visualize Decision Boundary
def plot_decision_boundary(X, y, model):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                          np.arange(y_min, y_max, 0.01))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Pa
    plt.title("K-NN Decision Boundary")
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 2")
    plt.show()

plot_decision_boundary(X_test, y_test, model)
```

```
# Evaluation Metrics
print("Classification Report:\n", classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=mode)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()
```




```
Classification Report:
              precision    recall  f1-score   support

     0       1.00      1.00      1.00        16
     1       1.00      1.00      1.00        14

 accuracy          1.00          1.00          1.00        30
 macro avg          1.00          1.00          1.00        30
weighted avg          1.00          1.00          1.00        30
```

Confusion Matrix





# Explanation

## 1. Dataset Creation:

- A synthetic dataset with two features and two classes is generated using `make_classification`.

## 2. Model Training:

- A `KNeighborsClassifier` with ( `k=5` ) (5 neighbors) is trained using the training data.

## 3. Decision Boundary Visualization:

- The decision boundary is visualized using `contourf`, showing the regions where each class is predicted based on the nearest neighbors.

## 4. Evaluation Metrics:

- `classification_report` provides precision, recall, F1-score, and accuracy.
- `confusion_matrix` and its visualization display correct and incorrect predictions for each class.

## Key Metrics in Classification

- **Precision:**  $\frac{\text{TP}}{\text{TP} + \text{FP}}$
- **Recall:**  $\frac{\text{TP}}{\text{TP} + \text{FN}}$
- **F1-Score:** Harmonic mean of precision and recall.
- **Accuracy:**  $\frac{\text{TP} + \text{TN}}{\text{Total}}$

This example provides a hands-on implementation of K-NN classification with a visual understanding of its decision-making process and performance evaluation.



### ✓ 3.4. Naive Bayes: Practical Example

This example demonstrates using a Naive Bayes classifier in Python for classification, visualizing the decision boundaries, and evaluating its performance with metrics.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, (

# Generate synthetic dataset
X, y = make_classification(n_samples=100, n_features=2, n_informative=

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.


# Naive Bayes Model
model = GaussianNB()
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Visualize Decision Boundary
def plot_decision_boundary(X, y, model):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                          np.arange(y_min, y_max, 0.01))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Pa
    plt.title("Naive Bayes Decision Boundary")
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 2")
    plt.show()


plot_decision_boundary(X_test, y_test, model)
```


```
# Evaluation Metrics
print("Classification Report:\n", classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=mode)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()
```



Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.94   | 0.97     | 16      |
| 1            | 0.93      | 1.00   | 0.97     | 14      |
| accuracy     |           |        | 0.97     | 30      |
| macro avg    | 0.97      | 0.97   | 0.97     | 30      |
| weighted avg | 0.97      | 0.97   | 0.97     | 30      |





# Explanation

## 1. Dataset Creation:

- A synthetic dataset with two features and two classes is generated using `make_classification`.

## 2. Model Training:

- A `GaussianNB` (Gaussian Naive Bayes) model is trained using the training data.

## 3. Decision Boundary Visualization:

- The decision boundary is visualized using `contourf`, showing the regions where each class is predicted based on the probabilistic calculations.

## 4. Evaluation Metrics:

- `classification_report` provides precision, recall, F1-score, and accuracy for each class.
- `confusion_matrix` and its visualization display correct and incorrect predictions for each class.

---

## Key Metrics in Classification

- **Precision:**  $\frac{\text{TP}}{\text{TP} + \text{FP}}$
- **Recall:**  $\frac{\text{TP}}{\text{TP} + \text{FN}}$
- **F1-Score:** Harmonic mean of precision and recall.
- **Accuracy:**  $\frac{\text{TP} + \text{TN}}{\text{Total}}$

This concise example demonstrates how Naive Bayes classifiers work in practice, with visual insights into their decision-making and performance evaluation.

### ✓ 3.5. Support Vector Machines (SVM): Practical Example

This example demonstrates using an SVM for classification in Python, visualizing the decision boundaries, and evaluating its performance.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix, (

# Generate synthetic dataset
X, y = make_classification(n_samples=100, n_features=2, n_informative=

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.


# SVM Model with a linear kernel
model = SVC(kernel='linear', probability=True, random_state=42)
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Visualize Decision Boundary
def plot_decision_boundary(X, y, model):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                          np.arange(y_min, y_max, 0.01))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.Paired)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.Pa
    plt.title("SVM Decision Boundary")
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 2")
    plt.show()

plot_decision_boundary(X_test, y_test, model)
```

```
# Evaluation Metrics
print("Classification Report:\n", classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=mode)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()
```




Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 16      |
| 1            | 1.00      | 1.00   | 1.00     | 14      |
| accuracy     |           |        | 1.00     | 30      |
| macro avg    | 1.00      | 1.00   | 1.00     | 30      |
| weighted avg | 1.00      | 1.00   | 1.00     | 30      |

**Confusion Matrix**





# Explanation

## 1. Dataset Creation:

- A synthetic dataset with two features and two classes is generated using `make_classification`.

## 2. Model Training:

- A `SVC` (Support Vector Classifier) with a linear kernel is trained to find the optimal hyperplane separating the classes.

## 3. Decision Boundary Visualization:

- The decision boundary is visualized using `contourf`, showing the regions classified for each class. Support vectors influence the boundary and are implicitly highlighted by the classifier.

## 4. Evaluation Metrics:

- `classification_report` provides precision, recall, F1-score, and accuracy for each class.
- `confusion_matrix` and its visualization display correct and incorrect predictions for each class.

# Key Metrics in Classification

- **Precision:**  $\frac{\text{TP}}{\text{TP} + \text{FP}}$
- **Recall:**  $\frac{\text{TP}}{\text{TP} + \text{FN}}$
- **F1-Score:** Harmonic mean of precision and recall.
- **Accuracy:**  $\frac{\text{TP} + \text{TN}}{\text{Total}}$

This example gives a hands-on implementation of SVMs with a focus on visualizing decision boundaries and evaluating the model's performance.



Here is a comparison table summarizing the key aspects of the classification algorithms you mentioned:

| Criteria                 | Logistic Regression                                    | Decision Trees                                  |
|--------------------------|--------------------------------------------------------|-------------------------------------------------|
| Type                     | Classification (Binary or Multi-class)                 | Classification (Binary or Multi-class)          |
| Model Type               | Linear model for classification                        | Non-linear, tree-based model                    |
| Decision Boundary        | Linear decision boundary                               | Non-linear decision boundary                    |
| Assumptions              | Assumes a linear relationship between input and output | Assumes data can be split by axis-aligned lines |
| Interpretability         | Easy to interpret with coefficients                    | Easy to interpret tree structure                |
| Computational Complexity | $O(n)$ for prediction (simple)                         | $O(\log(n))$ for prediction (balanced)          |
| Scalability              | Works well for small to medium-sized datasets          | Can struggle with large datasets                |
| Sensitivity to Noise     | Sensitive to outliers and irrelevant features          | Sensitive to noisy data (prone to overfitting)  |
| Regularization           | Yes, regularization via L1/L2 penalties                | Pruning can be used to prevent overfitting      |
| Advantages               | Simple, fast, interpretable                            | Handles both numerical and categorical data     |
| Disadvantages            | Assumes linearity, sensitive to outliers               | Prone to overfitting, complex boundaries        |

Key Takeaways:

- **Logistic Regression** is best for simple, linear relationships, and is fast and interpretable but assumes linearity.
- **Decision Trees** are non-linear and easy to interpret, but they are prone to overfitting and can be sensitive to noise.
- **K-NN** is non-parametric and flexible but can be slow for large datasets and sensitive to noise.
- **Naive Bayes** is efficient and works well for small datasets with independent features, but it may struggle with correlated features.
- **SVM** is effective for high-dimensional and complex datasets, especially with non-linear decision boundaries, but it can be computationally expensive and difficult to interpret.

Each algorithm has its own strengths and is suitable for different types of data and problems.

Here is a comparison table summarizing the different evaluation metrics used to assess the performance of classification models:

| Metric               | Accuracy                                                                                        |
|----------------------|-------------------------------------------------------------------------------------------------|
| <b>Definition</b>    | The proportion of correct predictions (both positive and negative) out of all predictions made. |
| <b>Formula</b>       | $\frac{TP + TN}{TP + TN + FP + FN}$                                                             |
| <b>Range</b>         | 0 to 1, where 1 is perfect accuracy.                                                            |
| <b>Best When</b>     | When classes are balanced and overall performance is important.                                 |
| <b>Disadvantages</b> | Can be misleading when the dataset is imbalanced, as it may favor the majority class.           |
| <b>Use Cases</b>     | Suitable for general model evaluation, especially when classes are balanced.                    |

## Key Takeaways:

- **Accuracy** is a good overall metric but may not be reliable with imbalanced datasets.
- **Precision** focuses on how many of the predicted positives are actually positive.
- **Recall (Sensitivity)** is concerned with how many of the actual positives are correctly predicted.
- **F1-Score** is a good metric when you need to balance Precision and Recall, especially in imbalanced datasets.
- **Confusion Matrix** provides a detailed breakdown of the classification performance, helping to understand the true positives, false positives, true negatives, and false negatives.

